

# On the CUDA implementation of Reduce and Scan

Cosmin Oancea and Troels Henriksen

September 2023

Course Contents

Implementation of Reduce


Implementation of Scan

## Summing an Integer Array

$$\sum_{i < n} x[i]$$

# Binary Tree Reduction

**The idea:** each thread reads two neighbouring elements, adds them together, and writes one element. This halves the array in size. Continue until only a single element is left.



- Each level becomes a kernel invocation, with number of threads equal to half the number of array elements.
- $O(n)$  work and  $O(\log(n))$  span (optimal).
- **Why is this not efficient?**

# Improving the Tree Reduction

**The idea:** instead of shrinking the array by a factor of two for each level, shrink it by the CUDA-block size.

- ▶ Same asymptotic performance.
- ▶ Avoids kernels with very few threads. E.g with block size 256:  
 $10000000 \rightarrow 39063 \rightarrow 153 \rightarrow 1$ .

# Improving the Tree Reduction

**The idea:** instead of shrinking the array by a factor of two for each level, shrink it by the CUDA-block size.

- ▶ Same asymptotic performance.
- ▶ Avoids kernels with very few threads. E.g with block size 256:  
 $10000000 \rightarrow 39063 \rightarrow 153 \rightarrow 1$ .

Implementation	$n = 1000$	$n = 1000000$
Tree reduction	$77\mu s$	$363\mu s$
Block reduction	$17\mu s$	$179\mu s$

# Applying Brent's Lemma

**The idea:** instead of letting the thread count depend on the input size, always launch the same number of threads, and have each thread perform an efficient sequential summation of a *chunk* of the input.

- ▶ GPUs have a maximum (hardware/problem-dependent) capacity for exploiting parallelism. Beyond that limit, parallelism is at best worthless, and usually comes with overhead (e.g. excessive synchronisation).
- ▶ *A straightforward implementation of this idea only works if the operator is commutative, and also allows fusing a map producer!*

# Applying Brent's Lemma

**The idea:** instead of letting the thread count depend on the input size, always launch the same number of threads, and have each thread perform an efficient sequential summation of a *chunk* of the input.

- ▶ GPUs have a maximum (hardware/problem-dependent) capacity for exploiting parallelism. Beyond that limit, parallelism is at best worthless, and usually comes with overhead (e.g. excessive synchronisation).
- ▶ *A straightforward implementation of this idea only works if the operator is commutative, and also allows fusing a map producer!*

Implementation	$n = 1000$	$n = 1000000$
Tree reduction	$77\mu s$	$363\mu s$
Block reduction	$17\mu s$	$179\mu s$
Chunked reduction	$70\mu s$	$103\mu s$



# Using Atomics

**The idea:** GPUs have special hardware support for performing certain memory updates atomically. In CUDA, this is exposed through *atomic operations*.

```
int atomicAdd( volatile __global int *p  
              , int val)
```

- ▶ Concise parallel reduction: each thread reads an element and uses `atomicAdd()` to update the same location in memory.
- ▶ **Why is this slow for large inputs?**

# Using Atomics

**The idea:** GPUs have special hardware support for performing certain memory updates atomically. In CUDA, this is exposed through *atomic operations*.

```
int atomicAdd( volatile __global int *p  
              , int val)
```

- Concise parallel reduction: each thread reads an element and uses `atomicAdd()` to update the same location in memory.
- **Why is this slow for large inputs?**

Implementation	$n = 1000$	$n = 1000000$
Tree reduction	$77\mu s$	$363\mu s$
Block reduction	$17\mu s$	$179\mu s$
Chunked reduction	$70\mu s$	$103\mu s$
Atomics	$8\mu s$	$1278\mu s$

# Coalesced Access to Global Memory

- ▶ On NVIDIA GPUs, the hardware threads are split into WARPS, where a WARP consists of 32 threads.
- ▶ The threads in a WARP execute in lockstep—in SIMD fashion—meaning at each given cycle they all execute the same instruction.
- ▶ Coalesced access to global memory is obtained when the threads in a WARP access in their common load/store instruction consecutive words in global memory.
- ▶ A coalesced load/store instruction is serviced by one (possibly two) memory transactions, hence fast.
- ▶ Uncoalesced access may requires as many as 32 memory transactions, which are sequentially issued by the memory controller, hence may generate significant slowdown.
- ▶ Multi-threading somehow alleviates the “uncoalesced” overhead but only to some extent.

# For Non-Commutative Operators (MSSP)

We denote the CUDA block size with  $B$ .

- ▶ We can also fuse the mapped function in the reduction itself.
- ▶ Have each thread reducing sequentially some CHUNK **consecutive** elements (statically known):
  - ▶ then the  $B$  per-thread results are reduced cooperatively by the threads in a block.
  - ▶ sequential (virtualization) loop on top so that the number of blocks is kept to under 1024, i.e., two-stage reduce.
- ▶ However, if done naively, this will lead to **uncoalesced** access. For **coalesced**, use shared memory as a staging buffer:
  - ▶ the threads in a block read consecutive elements from global memory and write them in shared memory;
  - ▶ then each thread processes sequentially its CHUNK consecutive elements from shared memory, which is not subject to the “coalescing” overhead.

Course Contents


Implementation of Reduce

Implementation of Scan

# Parallel Exclusive Scan with Associative Operator $\oplus$

Two Steps:

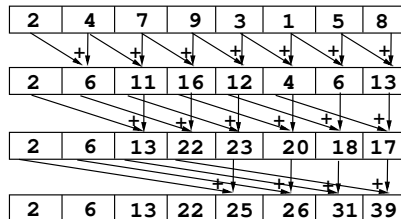
- ▶ **Up-Sweep**: similar with reduction
- ▶ Root is replaced with neutral element.
- ▶ **Down-Sweep**:
  - ▶ the left child sends its value to parent and updates its value to that of parent.
  - ▶ the right-child value is given by  $\oplus$  applied to the left-child value and the (old) value of parent.
  - ▶ note that the right child is in fact the parent, i.e., in-place algorithm.



*Up-Sweep & Down-Sweep*

Scan's Work and Depth:  $D(n) = \Theta(\lg n)$ ,  $W(n) = \Theta(n)$

# Warp-Level Inclusive Scan for GPUs



Input: array A of  $n=2^k$  elements  
of type  $\alpha$

$\oplus : (\alpha, \alpha) \rightarrow \alpha$  associative

Output:  $B = [a_1, a_1 \oplus a_2, \dots, \oplus_{j=0}^{n-1} a_j]$

1. forall  $i = 0 : n-1$  do
2.      $B[i] \leftarrow A[i]$
3. endfor
4. for  $d = 0$  to  $k-1$  do
5.      $h = 2^d$
6.     forall  $i = h$  to  $n-1$  do
7.          $B[i] \leftarrow B[i-h] \oplus B[i]$
8.     endfor
9. endfor

Offers better performance because it operates in one sweep rather than two!

# Warp-Level Inclusive Scan Implementation

Input: array A of  $n=2^k$  elements  
of type  $\alpha$

$\oplus : (\alpha, \alpha) \rightarrow \alpha$  associative

Output:  $B = [a_1, a_1 \oplus a_2, \dots, \oplus_{j=0}^{n-1} a_j]$

1. forall  $i = 0 : n-1$  do

2.  $B[i] \leftarrow A[i]$

3. endfor

4. for  $d = 0$  to  $k-1$  do


5.  $h = 2^d$

6. forall  $i = h$  to  $n-1$  do

7.  $B[i] \leftarrow B[i-h] \oplus B[i]$

8. endfor

9. endfor



- ▶ Open file "pbbKernels.cu.h", and implement function named "scanIncWarp" (follow the instructions)
- ▶ Your  $n = \text{WARP}$  and  $k = \lg \text{WARP}$ ; ignore the init loop;
- ▶ Unroll the for  $d$  loop (#pragma unroll);
- ▶ loop forall  $i = h$  to  $n-1$  is implicit (parallel threads),
- ▶ it should be replaced by a condition if  $(i \geq h) \{ \dots \}$ ,
- ▶ except that  $i$  in condition if  $(i \geq h)$  is **not the thread id**.
- ▶ Remember, you want to scan each warp, independently!



# OpenCL Scan Implementation

## BLACKBOARD!

- ▶ Generic CPU skeleton in `hostSkel.cu.h`.
  1. reduce each block and publish the per-block results in a buffer `buff` (number of blocks  $\leq 1024$ )
  2. scan in place the buffer `buff` using one CUDA block.
  3. now you can implement the core scan kernel:
    - 3.1 each thread reads and scans its CHUNK of consecutive elements held in register memory; (The elements are copied in a coalesced way by using shared memory as a staging buffer.)
    - 3.2 the last element (reduction) of each thread's chunk (held in registers) is published in shared memory.
    - 3.3 then the threads in the block cooperatively scan the per thread reductions.
    - 3.4 then each thread `tid` updates the elements of its chunk array with the element at position `tid-1` from step 3.3, and with the previous-block element of `buff` from step 2.
    - 3.5 the CHUNK per-thread results are copied from register to shared to global memory (to ensure coalesced writes)
- ▶ If  $N$  is the length of the input, this requires  $2N$  coalesced reads and  $N$  coalesced writes from/to global memory.